

## Python Candidate Test – RAG Pipeline with Streamlit + LangChain

### 🎯 Objective

Build a Retrieval-Augmented Generation (RAG) chatbot using Python that:  
Ingests custom .txt content (e.g., blog posts)

Embeds and indexes it using FAISS

Uses LangChain to retrieve relevant chunks

Uses a Hugging Face model (GPT2 or Mistral) to answer queries

Includes a basic Streamlit UI for interaction

### ✓ Deliverables

A Colab-compatible .ipynb notebook

A Streamlit app script (app.py) to run the chatbot locally

Sample output from 2 questions

Link to GitHub repo or Google Drive (if required)

### 📁 Project Structure

```
rag_test/
├── blogs/
│   ├── blog1.txt
│   └── blog2.txt
├── rag_pipeline.ipynb    ✓ (main notebook for logic)
├── app.py                ✓ (streamlit UI)
└── requirements.txt      ✓ (dependencies)
```

### 🛠 Tools to Use

LangChain (for pipeline orchestration)

sentence-transformers (all-MiniLM-L6-v2 or bge-small-en)

faiss (vector store)

transformers (for LLM)

Streamlit (for frontend UI)

### 🔑 Task Breakdown

✓ Step 1: Data Ingestion & Preprocessing

Load text from files in the blogs/ folder.

Chunk text into ~200-word sections with titles preserved.

Clean basic punctuation and newlines.

#### ✓ Step 2: Embedding + Vector Store

Use sentence-transformers to generate embeddings.

Store in FAISS with document metadata.

Save and reload FAISS index for persistence.

#### ✓ Step 3: LangChain Integration

Create a FAISS Retriever with LangChain

Use ConversationalRetrievalChain with:

A basic prompt template

A LLM like GPT2 or any Hugging Face small model

Allow queries like:

"What is energy mastery?"

"How do high performers avoid burnout?"

#### ✓ Step 4: Streamlit UI

Create a simple UI with:

Title: "Your Personal Knowledge Chatbot"

Textbox for user input

Chat history below

Call backend via LangChain pipeline

Display model's answers cleanly

```
import streamlit as st
```

```
st.set_page_config(page_title="Knowledge Chatbot")
```

```
st.title("📖 Your Personal Knowledge Chatbot")
```

```
if "history" not in st.session_state:  
    st.session_state.history = []
```

```
query = st.text_input("Ask something from your content:")
```

```
if query:
    from backend import get_answer # optional helper

    response = get_answer(query)
    st.session_state.history.append((query, response))
```

```
for q, r in reversed(st.session_state.history):
    st.markdown(f"***You:** {q}")
    st.markdown(f"***Bot:** {r}")
```

🔗 Sample Prompt Template (LangChain)

template = """You are a helpful assistant. Use the context below to answer the user's question.

Context:  
{context}

Question:  
{question}

Answer: """

📄 requirements.txt  
transformers  
sentence-transformers  
faiss-cpu  
langchain  
streamlit

▣ Sample Test Data: blog1.txt

link

✓ Evaluation Criteria

Area

Points

Clean Python structure & documentation

10

FAISS + embedding setup

10

LangChain integration

10

Working Streamlit UI

10

Answer relevance from model

10

Total

50

📁 Bonus (Optional)

Add conversation memory

Replace GPT2 with Mistral (via Ollama) or Mixtral (if Colab supports it)

Use ChromaDB instead of FAISS